

In [5]:

```
import numpy as np
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
import matplotlib.axes as ax
from IPython.display import clear_output
```

In [6]:

```
df = pd.read_csv("data_for_lr.csv")
df.head()

# Drop null values
df = df.dropna()
df.head()
```

Out[6]:


	x	y
0	24.0	21.549452
1	50.0	47.464463
2	15.0	17.218656
3	38.0	36.586398
4	87.0	87.288984

In [7]:

```
sns.pairplot(
    df[:100]
)
```

Out[7]:

<seaborn.axisgrid.PairGrid at 0x161e928ed90>



Train-Test Split

In [8]:

```
# skipping cross-validation part for now for convenience
N = len(df)
training_data_len = int (N*0.8)
x_train, y_train = np.array(df['x'][0:training_data_len]).reshape(training_data_len, 1),
np.array(df['y'][0:training_data_len]).reshape(
    training_data_len, 1
)
x_test, y_test = np.array(df['x'][training_data_len:N]).reshape(N - training_data_len, 1),
np.array(
    df['y'][training_data_len:N]).reshape(N - training_data_len, 1)
```

In [9]:

N

Out[9]:

699

linear regression

here's the mathematical derivation I did on paper before implementing it in python

SUHANI CHAWLA 2/13/2024

Discriminative model \leftrightarrow generative model

$P(y|x, \theta) \propto \underbrace{P(y)}_{\text{prior}} \underbrace{P(x|y, \theta)}_{\text{likelihood}}$

$y = \sum \theta_i x_i + \epsilon$ $\epsilon \sim N(0, \sigma)$ $\epsilon \sim \text{noise}$

$P(\epsilon) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{\epsilon^2}{2\sigma^2}}$

$P(y|x, \theta) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y^{(1)} - \sum \theta x^{(1)})^2}{2\sigma^2}}$

$L(\theta) = \prod P(y_1|x_1, \theta) P(y_2|x_2, \theta) \dots$

$L(\theta) = \prod_{i=1}^n P(y^{(i)}|x^{(i)}, \theta)$

$\log \text{ likelihood } L(\theta) = \sum_{i=1}^n \log(P(y^{(i)}|x^{(i)}, \theta)) = \sum_{i=1}^n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} \sum (y^i - \theta x^i)^2$

To minimize $L(\theta)$, ~~we want~~ ^{try to} minimize $\sum (y^i - \theta x^i)^2$

Minimize $J(\theta) = \frac{1}{2} \sum (y^i - \theta x^i)^2$ ^{square error}

$$J(\theta) = \frac{1}{2} \sum (x\theta - y)^2$$

$$\frac{d}{d\theta} J(\theta) = x^T x \theta - x^T y = 0 \rightarrow \text{for minimizing}$$

$$\theta = (x^T x)^{-1} x^T y$$

$$\theta_J^t = \theta_J^{t-1} - \alpha \frac{d}{d\theta} J(\theta)$$

$$\theta_J^t = \theta_J^{t-1} - \alpha \sum_{i=1}^m (x^i \theta - y^i) x_J$$

~~learning rate~~ learning rate

Train \rightarrow Validation data \rightarrow Test data
stop on validation error

Shuffle

$$SST = \sum_{i=1}^n (y^i - \bar{y})^2$$

$$SSR = \sum (x^i \theta - y^i)^2$$

$$R^2 \text{ score} = \frac{SSR}{SST} = 1 - \frac{SSR}{SST}$$

higher the R^2 score, better the model

$$\text{mean square error} = \frac{1}{n} \sum_{i=1}^n \frac{1}{2} (x\theta - y)^2$$

In [10]:

```
# the same calculation in python
```

```
class LinearRegression:
```

```
    def __init__(self):
```

```
        self.Q0 = np.random.uniform(0, 1) * -1 # Intercept
```

```
        self.Q1 = np.random.uniform(0, 1) * -1 # Coefficient of X
```

```
        self.losses = [] # Storing the loss of each iteration
```

```
        self.r2s = []
```

```
    def forward_propogation(self, training_input):
```

```
        predicted_values = np.multiply(self.Q1, training_input) + self.Q0 #  $y = mx + c$ 
```

```
        return predicted_values
```

```
    def cost(self, predictions, training_output):
```

```
        lmse = np.mean((predictions - training_output) ** 2) # Calculating the cost
```

```
        ssr = np.sum((predictions - training_output)**2)
```

```
        sst = np.sum((training_output - np.mean(training_output))**2)
```

```
        r2_score = 1 - (ssr/sst)
```

```
        return lmse, r2_score
```

```
    def finding_derivatives(self, cost, predictions, training_input, training_output):
```

```
        diff = predictions - training_output
```

```
        dQ0 = np.mean(diff) #  $d(J(Q0, Q1))/d(Q0)$ 
```

```
        dQ1 = np.mean(np.multiply(diff, training_input)) #  $d(J(Q0, Q1))/d(Q1)$ 
```

```
        return dQ0, dQ1
```

```
    def train(self, x_train, y_train, lr, itrs):
```

```
        for i in range(itrs):
```

```
            # Finding the predicted values (Using the linear equation  $y=mx+c$ )
```

```
            predicted_values = self.forward_propogation(x_train)
```

```
            # Calculating the Loss
```

```
            loss, r2 = self.cost(predicted_values, y_train)
```

```
            self.losses.append(loss)
```

```
            self.r2s.append(r2)
```

```
            # Back Propagation (Finding Derivatives of Weights)
```

```
            dQ0, dQ1 = self.finding_derivatives(
```

```
                loss, predicted_values, x_train, y_train
```

```
            )
```

```
            # Updating the Weights
```

```
            self.Q0 = self.Q0 - lr * (dQ0)
```

```
            self.Q1 = self.Q1 - lr * (dQ1)
```

```
            # It will dynamically update the plot of the straight line
```

```
            line = self.Q0 + x_train * self.Q1
```

```
            clear_output(wait=True)
```

```
            plt.plot(x_train, y_train, "+", label="Actual values")
```

```
            plt.plot(x_train, line, label="Linear Equation")
```

```
            plt.xlabel("Train-X")
```

```
            plt.ylabel("Train-Y")
```

```
            plt.legend()
```

```
            plt.show()
```


```
        return (
```

```
            self.Q0,
```

```
        self.Q1,  
        self.losses,  
        self.r2s  
    ) # Returning the final model weights and the losses
```


In [11]:

```
lr = 0.0001 # Learning Rate  
itrs = 30 # No. of iterations  
model = LinearRegression()  
Q0, Q1, losses, r2s = model.train(x_train, y_train, lr, itrs)
```



In [12]:

```
plt.plot(losses)  
plt.show()  
print("from this graph we can see that after 10 iterations the loss is stabilized, therefore we could stop roughly after 10 iterations")
```




from this graph we can see that after 10 iterations the loss is stabilized, therefore we

could stop roughly after 10 iterations

In [13]:

```
plt.plot(r2s)
plt.show()
```



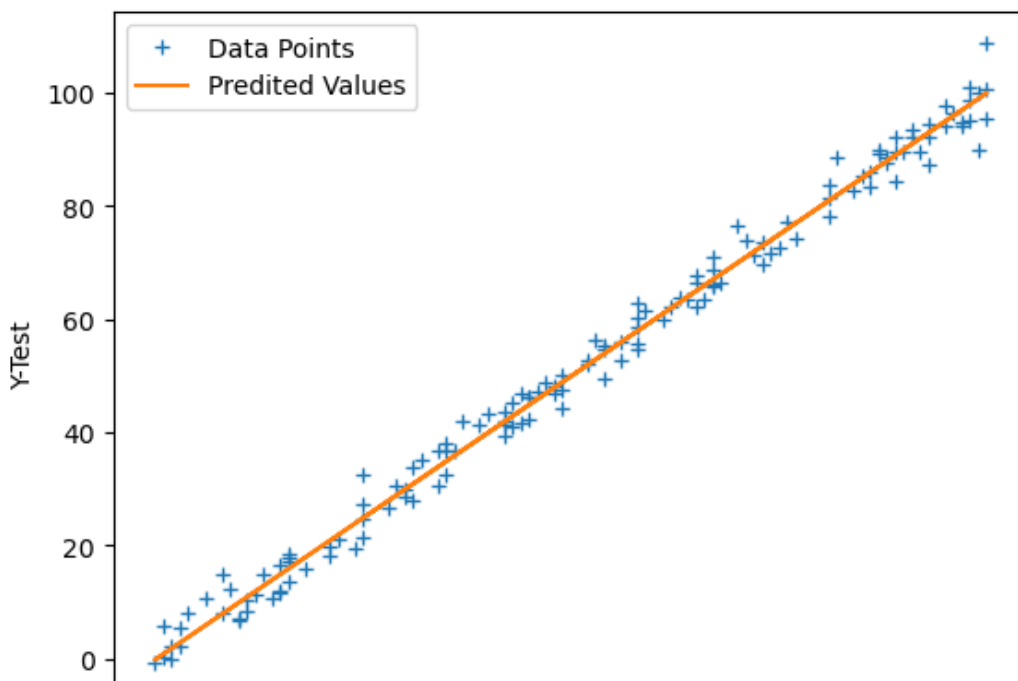
performance on test data


In [14]:

```
y_pred = model.forward_propagation(x_test)
```

In [15]:

```
plt.plot(x_test, y_test, "+", label="Data Points")
plt.plot(x_test, y_pred, label="Predited Values")
plt.xlabel("X-Test")
plt.ylabel("Y-Test")
plt.legend()
plt.show()
```





2nd way

the manual implementation was for my own understanding, now using sklearn LinearRegression and train_test_split

In [16]:

```
df2 = pd.read_csv("data_for_lr.csv")
df2.head()
```

```
# Drop null values
df2 = df2.dropna()
```

```
Y = df2['y']
X = df2.drop("y", axis=1)
len(X)==len(Y)
```

Out[16]:

True

In [17]:

```
X.head()
```

Out[17]:

	x
0	24.0
1	50.0
2	15.0
3	38.0
4	87.0

In [18]:

```
Y.head()
```

Out[18]:

```
0    21.549452
1    47.464463
2    17.218656
3    36.586398
4    87.288984
Name: y, dtype: float64
```

In [21]:

```
### do train test split using sklearn
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.25)
```


In [24]:

```
### fit a regression model on train data
### predict values of on test data
from sklearn.linear_model import LinearRegression
regr = LinearRegression()
regr.fit(X_train, Y_train)
print(regr.score(X_test, Y_test))
```

0.9900697525613219

In [42]:

```
# graphing results
import matplotlib.pyplot as plt
y_pred = regr.predict(X_test)
plt.scatter(X_test, Y_test, label = 'data_points')
plt.scatter(X_test, y_pred, color = 'r', label = 'predicted values')
plt.legend()
plt.show()
```



In []:

In []: